

# Computer Networks

## @CS.NYTU

### Lecture 2: Applications

Instructor: Kate Ching-Ju Lin (林靖茹)

Slides modified from

“Computer Networking: A Top-Down Approach” 7th Edition


# Outline

---

- Principles of network applications
- Web and HTTP
- E-Mail and SMTP
- DNS
- **Peer-to-peer applications**
- Video streaming and CDN

# Pure P2P Architecture


- No always-on server
  - Might have some **super nodes**, mainly used for management
- Arbitrary end systems (peers) directly communicate
  - **Self-scalable**
- Peers are intermittently connected and change IP addresses dynamically
- Examples:
  - File distribution (BitTorrent)
  - Streaming (PPstream)
  - VoIP (Skype)



# P2P: Challenges

---

- How to locate your peer & find what you want?
- Need some kind of “directory” or “look-up” service



- centralized
- distributed, using a hierarchal structure
- distributed, using a flat structure
- distributed, without structure (“flooding”)
- distributed, using “hybrid” structured/unstructured

# P2P: Challenges

---

- **Technical**

- **Scale**: up to hundred of thousands or millions of machines
- **Dynamics**: machines can come and go any time

- **Social, economic and legal**

- **Incentive** Issues: free-loader problem
- Vast majority of users are **free-riders**
- Most share no files and answer no queries
- A few individuals contributing to the “public good”
- **Copyrighted** content and privacy
- Trust & **security** issues

# BitTorrent: Popular P2P App

---

- Designed for large file (e.g., video) distribution
- Focused on *efficient fetching*, not search
  - Distribute same file to many peers
  - Single publisher, many downloaders
- Divide large file into many pieces (chunks)
  - Replicate different pieces on different peers
  - A peer with a complete piece can trade with others
  - Peer can (hopefully) assemble the entire file
- Allows simultaneous downloading
  - Retrieve different pieces from different peers simultaneously
- Usually need to prevent “free loading”

# BitTorrent Components

---

- **Seed**

- Peer with entire file
- Fragmented in pieces

- **Leacher**

- Peer with an incomplete copy of the file

- **Torrent file**

- Passive component
- Store summaries of the pieces to allow peers to verify their integrity

- **Tracker**

- Allows peers to find each other
- Returns a list of random peers

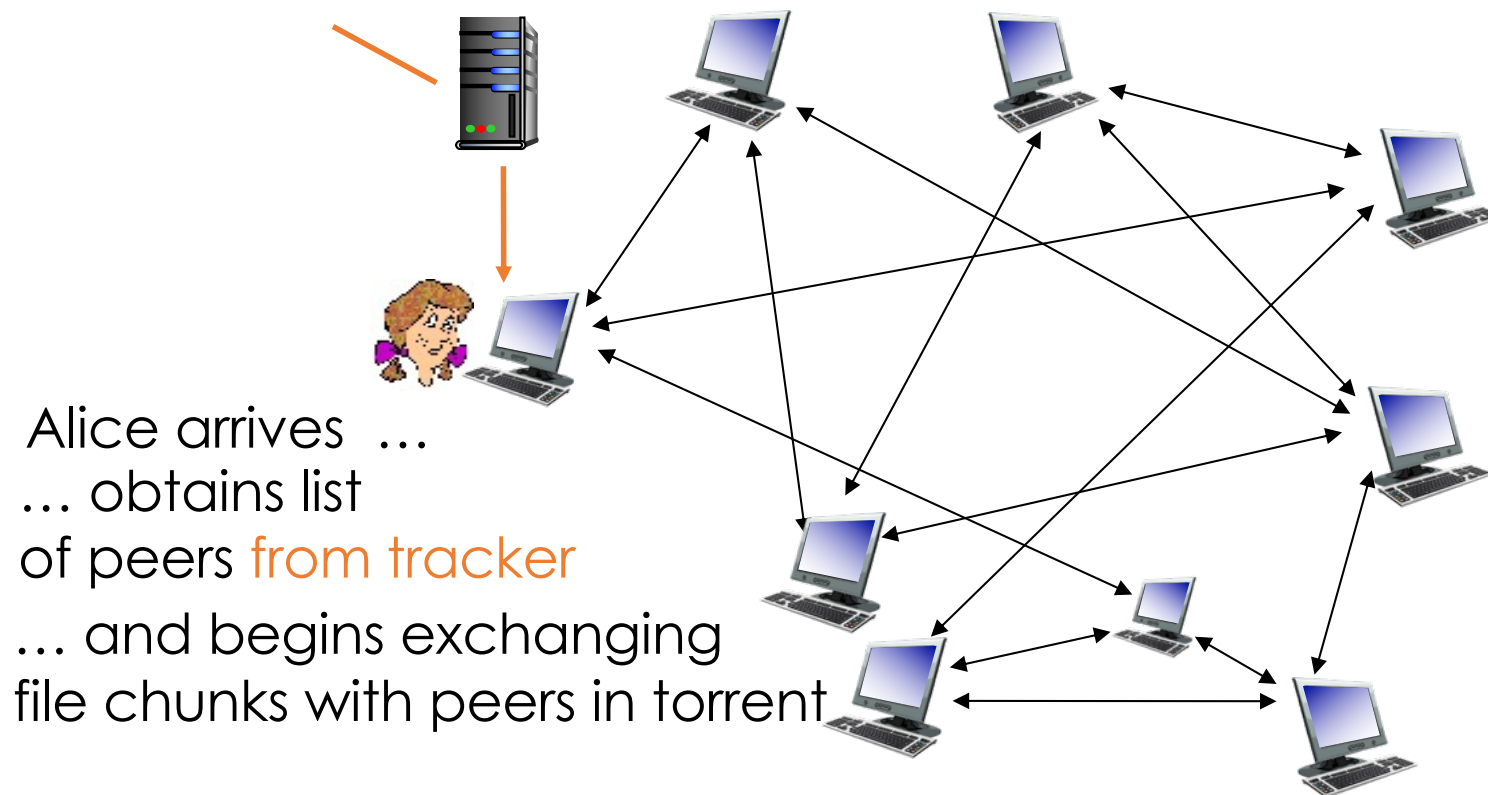
# BitTorrent Protocol

---

- File divided into 256Kb chunks
- Peers in torrent send/receive file chunks

**tracker:** tracks peers participating in torrent

**torrent:** group of peers exchanging chunks of a file





# BitTorrent Protocol


---

- Peer joining torrent:
  - Has no chunks, but will accumulate them over time from other peers
  - Register with tracker to get a list of peers (“neighbors”)
  - Exchange chunks with neighbors
- Peer may change neighbors over time
  - Tracker periodically suggests some random peers
  - Keep those random peers as neighbors if they are better (in terms of bandwidth or content)
- **Churn**: peers may come and go
- Once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent

# Bittorrent: Tit-for-Tat

---

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



# BitTorrent: Scheduling

---

Two choices should be made

- **Which chunk to download?**

- **rarest first**: download the chunk with the least copies first
- Why? More precious and might become disappear in the network
- Tracker should provide such information

- **Where to download?**

- **Highest rate first**: from a neighbor who can provide a higher downloading rate
- Update the neighbor list if the random peer suggested by the tracker supports a higher rate

# BT: requesting, sending file chunks

---

## *requesting chunks:*

- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first

## *sending chunks: tit-for-tat*

- Alice sends chunks to those peers currently sending her chunks at highest rate
  - other peers do not receive chunks from Alice
  - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
  - “optimistically unchoke” this peer
  - newly chosen peer may join top 4

# Outline

---

- Principles of network applications
- Web and HTTP
- E-Mail and SMTP
- DNS
- Peer-to-peer applications
- **Video streaming and CDN**

# “Multimedia” Networking

---

text

CS, NCTU  
Multimedia  
Networking  
2017 Spring  
ker ker

image



video




gaming



VR



audio / music




# Why Video Streaming Important?

---

- **High demand**

- Over 50% Internet traffic



- **High bandwidth requirement**

- 30, 60, 120 frames per seconds
- Video rate: 100kbps – 3Mbps or even >10Mbps (4K)


- **Complex compression algorithms**

- **No need to receive all the bits**
- Video quality is related to packet loss rate
  - But not linear proportional

# 1. Video-On-Demand Streaming

---

- Request **on demand**
- Online downloading and playing



**MOD**  
中華電信



**You** **Tube**

**NETFLIX**



## 2. Real-Time Video Streaming

- Record and deliver live video immediately
- Usually broadcasting



# 3. HTTP Streaming

---



- **DASH**: Dynamic, Adaptive Streaming over HTTP
- server:
  - Divide video file into multiple chunks
  - Each chunk stored, encoded at different rates
  - **Manifest file**: provides URLs for different chunks/rates
- client:
  - Periodically measures server-to-client bandwidth
  - Consulting manifest, requests one chunk at a time
    - Choose maximum coding rate sustainable given current bandwidth
  - Can choose different coding rates at different points in time (depending on available bandwidth at time)

# Why HTTP Streaming

---

- Traditional streaming delivered over RTP/UDP
- However, in today's Internet, content objects are stored **Content Delivery Networks (CDN)**
  - Many CDNs do not support RTP streaming
  - RTP often does not allow traffic through firewall
  - Each RTP stream is a separate session → not scalable
- Benefits of HTTP streaming
  - Firewall friendly
  - No need to maintain the session state on the server
  - Has been standardized as "ISO/IEC 23009-1, also known as MPEG-DASH" in Apr. 2012

# Why CDNs?

---

- **Challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?
- **Option 1:** single, large “mega-server”
  - Reliability: single point of failure
  - Not enough bandwidth: point of network congestion
  - Far from users: long path to distant clients
  - multiple copies of video sent over outgoing link

.... quite simply: this solution *doesn't scale*

# Why CDNs?

---

- **Challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?
- **Option 2:** store/serve multiple copies of videos at geographically distributed sites (CDN)
  - **enter deep:** push CDN servers deep into many access networks
    - ✓ close to users
    - ✓ used by Akamai, 1700 locations
  - **bring home:** smaller number (10's) of larger clusters in POPs near (but not within) access networks
    - ✓ used by Limelight


# Content Distribution Networks

---

- **CDN**: an application overlay (e.g., Akamai)
- Design Space
  - **Caching** (data-driven, passive)
    - explicit
    - transparent (hijacking connections)
  - **Replication** (pro-active)
    - server farms
    - geographically dispersed (CDN)
- Three Main CDN Providers (in North America, Europe):
  - Akamai, Limelight, Level 3 CDN

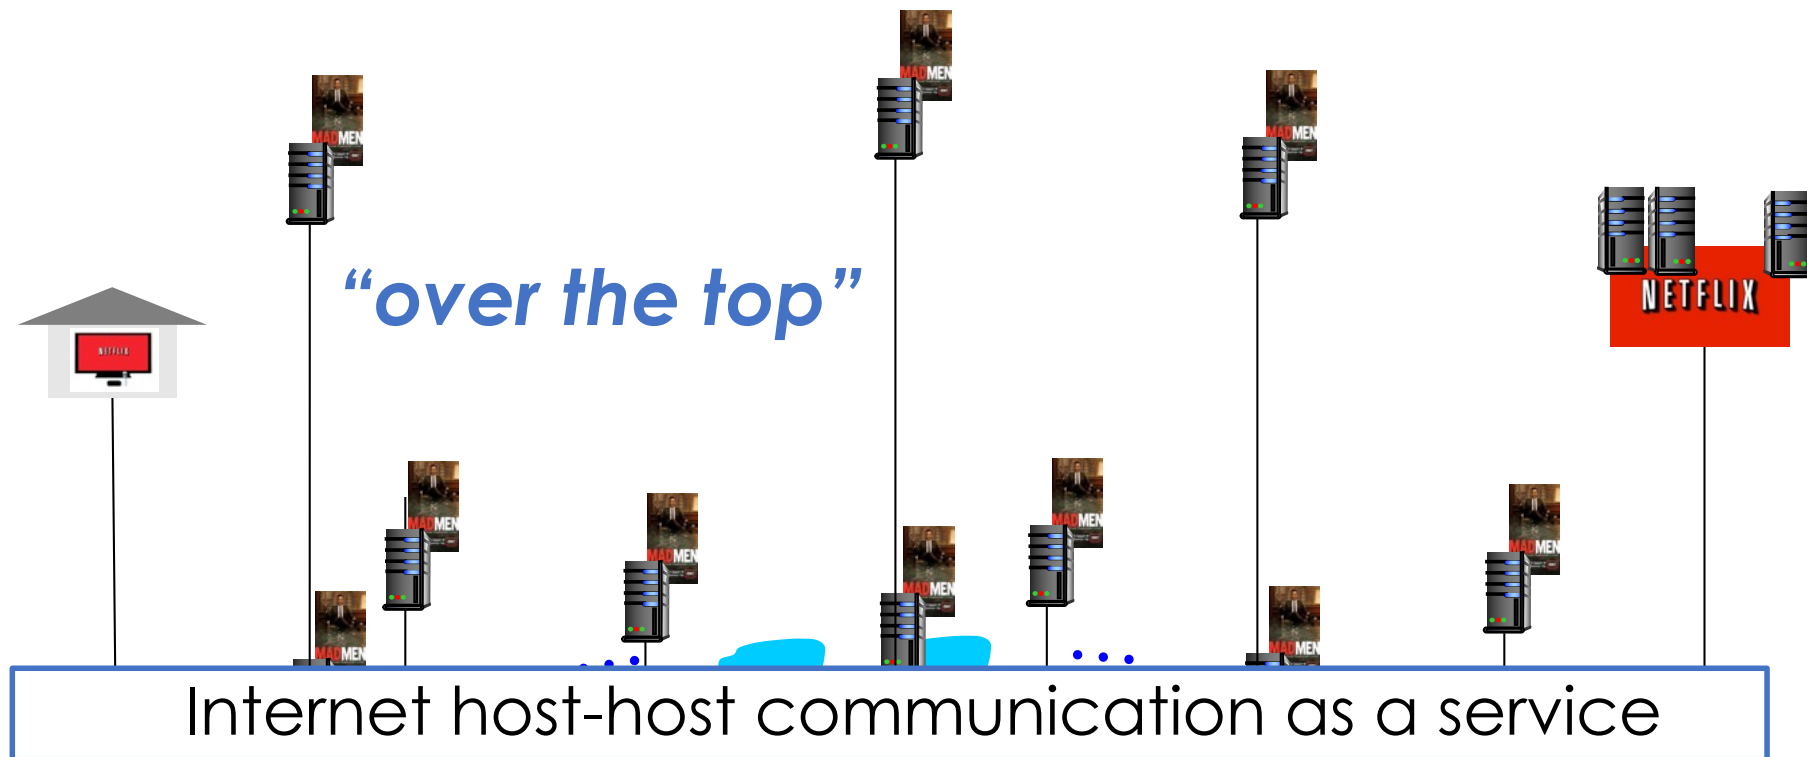
# Key Idea of CDN

- CDN: stores copies of content at CDN nodes
  - e.g. Netflix stores copies of *MadMen*
- Subscriber requests content from CDN
  - directed to nearby copy, retrieves content
  - may choose different copy if network path congested



# Framework of CDN

---



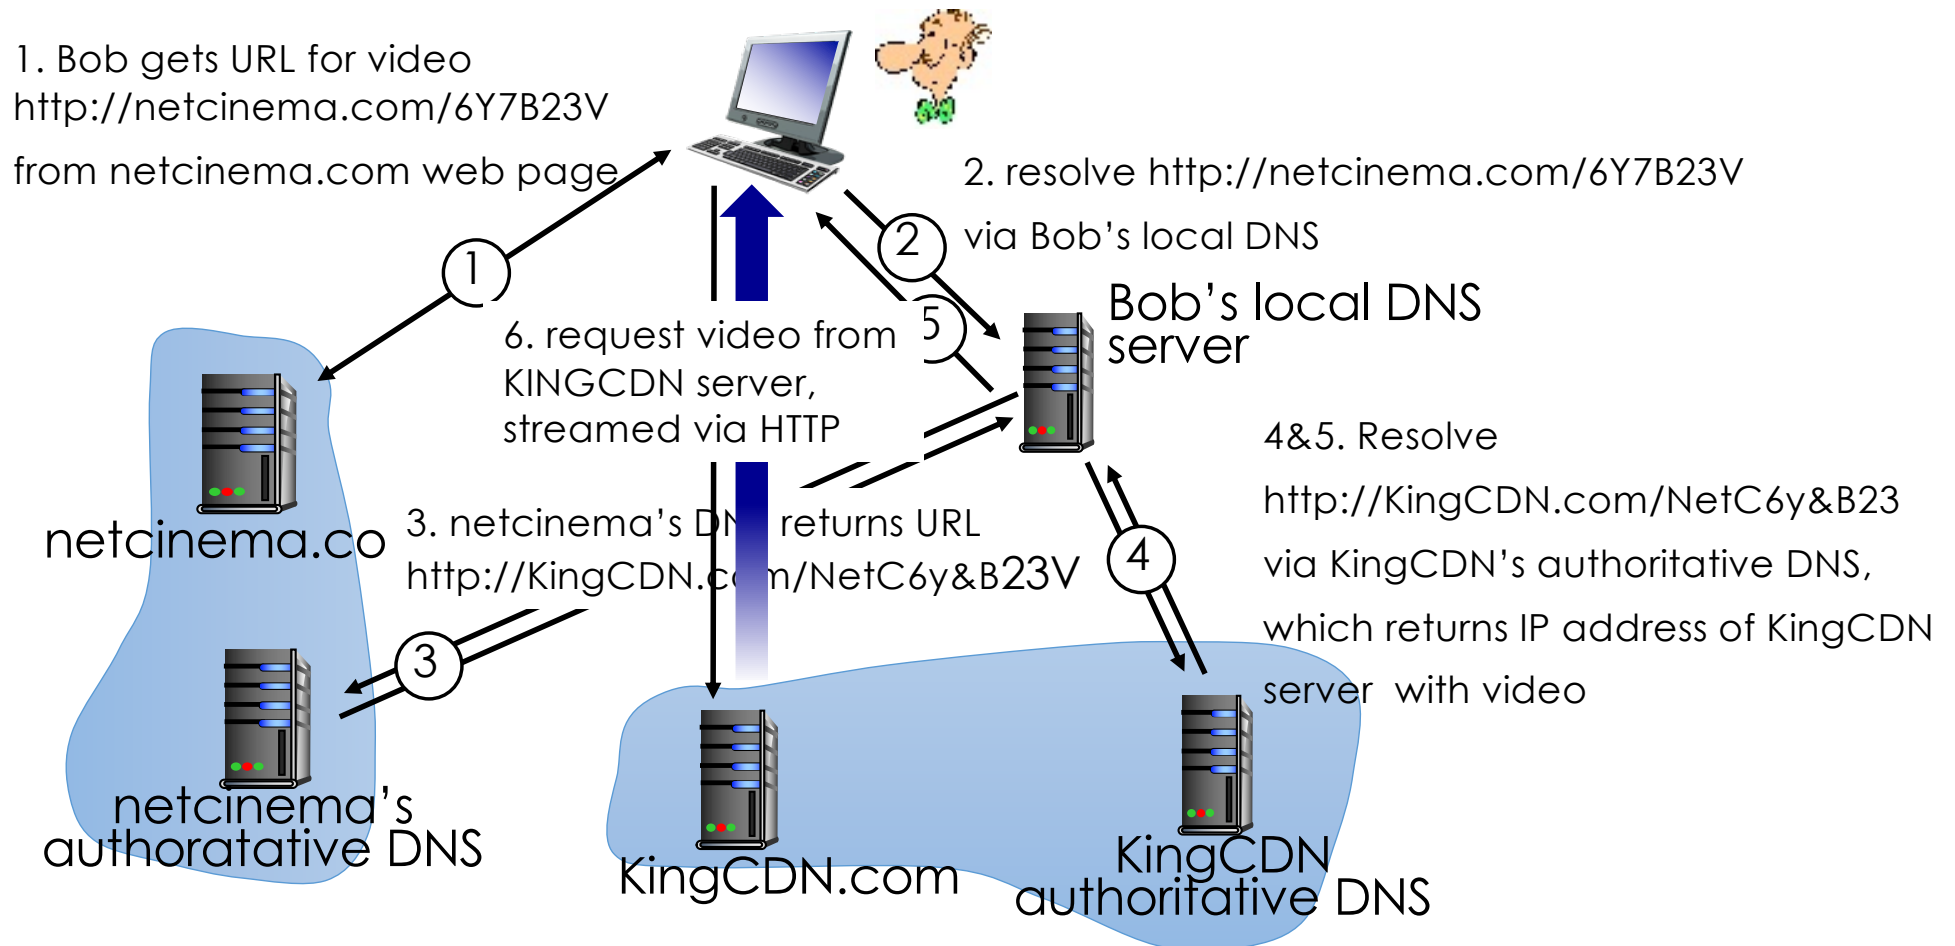
*OTT challenges:* coping with a congested Internet

- from which CDN node to retrieve content?
- viewer behavior in presence of congestion?
- what content to place in which CDN node?




# CDN Content Access

- Bob (client) requests video <http://netcinema.com/6Y7B23V>
- video stored in CDN at <http://KingCDN.com/NetC6y&B23V>



# Case Study: Netflix



# Cluster Selection Strategies

---

- **Geographically closest**

- Geographical distance  $\neq$  path length
- Local DNS location  $\neq$  user location
- Do not explicitly consider bandwidth and delay (varying with the congestion level)

- **Fastest response time**

- Periodically probe from CDN clusters to DNS servers
- Ingest huge probing traffic
- Some DNS servers configured to “not reply”